



IDA-ML I: Project Presentation

Project: Binary Stroke Classification

Examinee: Paul Utsch

Date: August 7, 2024

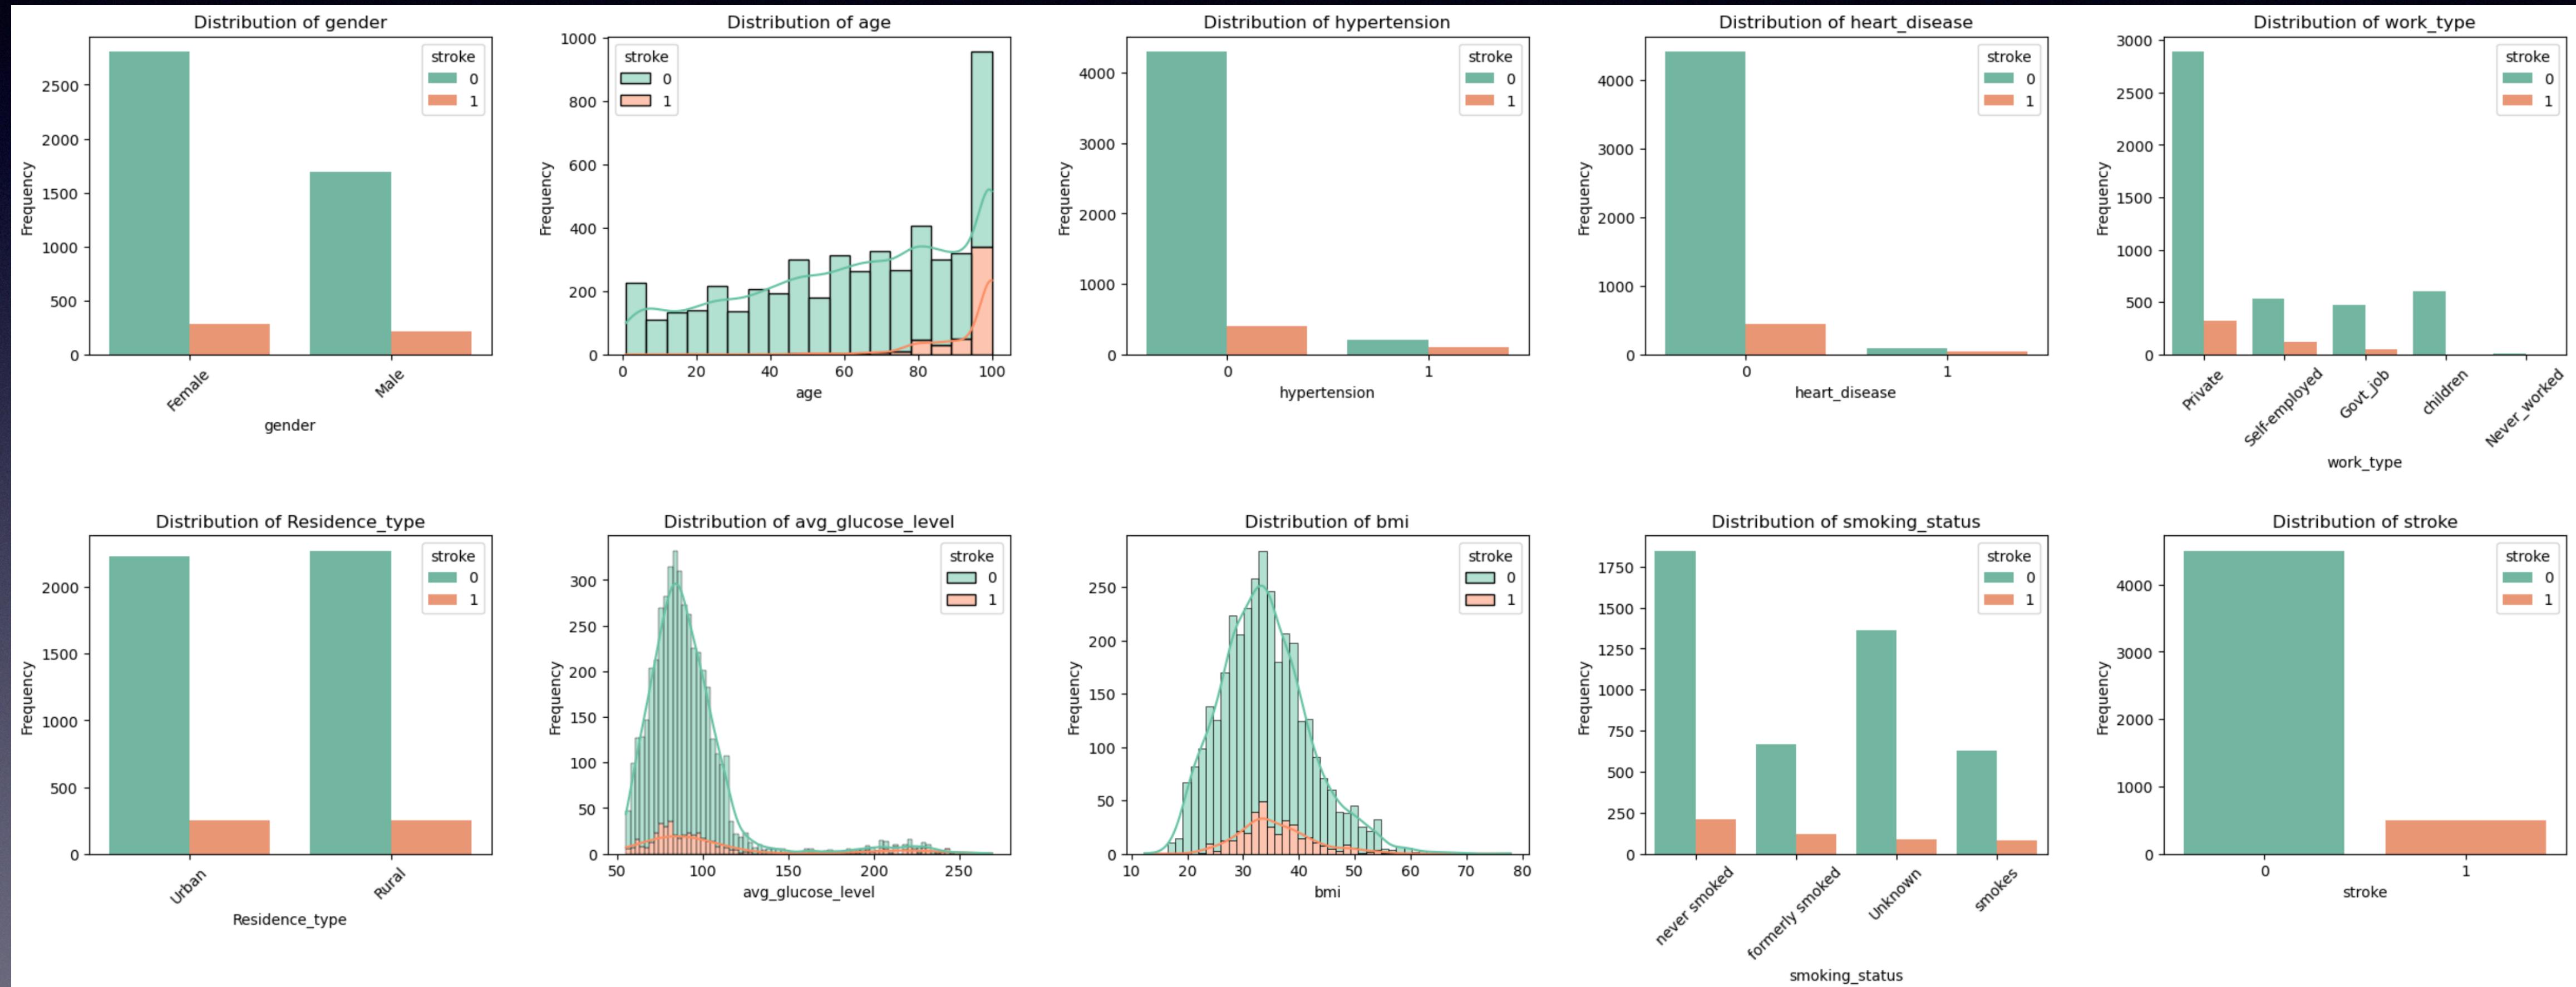
Table of Contents

1. Problem Setting
2. Data Analysis & Preprocessing
3. Machine Learning Methods
4. Feature Selection for Logistic Regression
5. Model Selection & Evaluation Protocol
6. Results
7. Conclusion

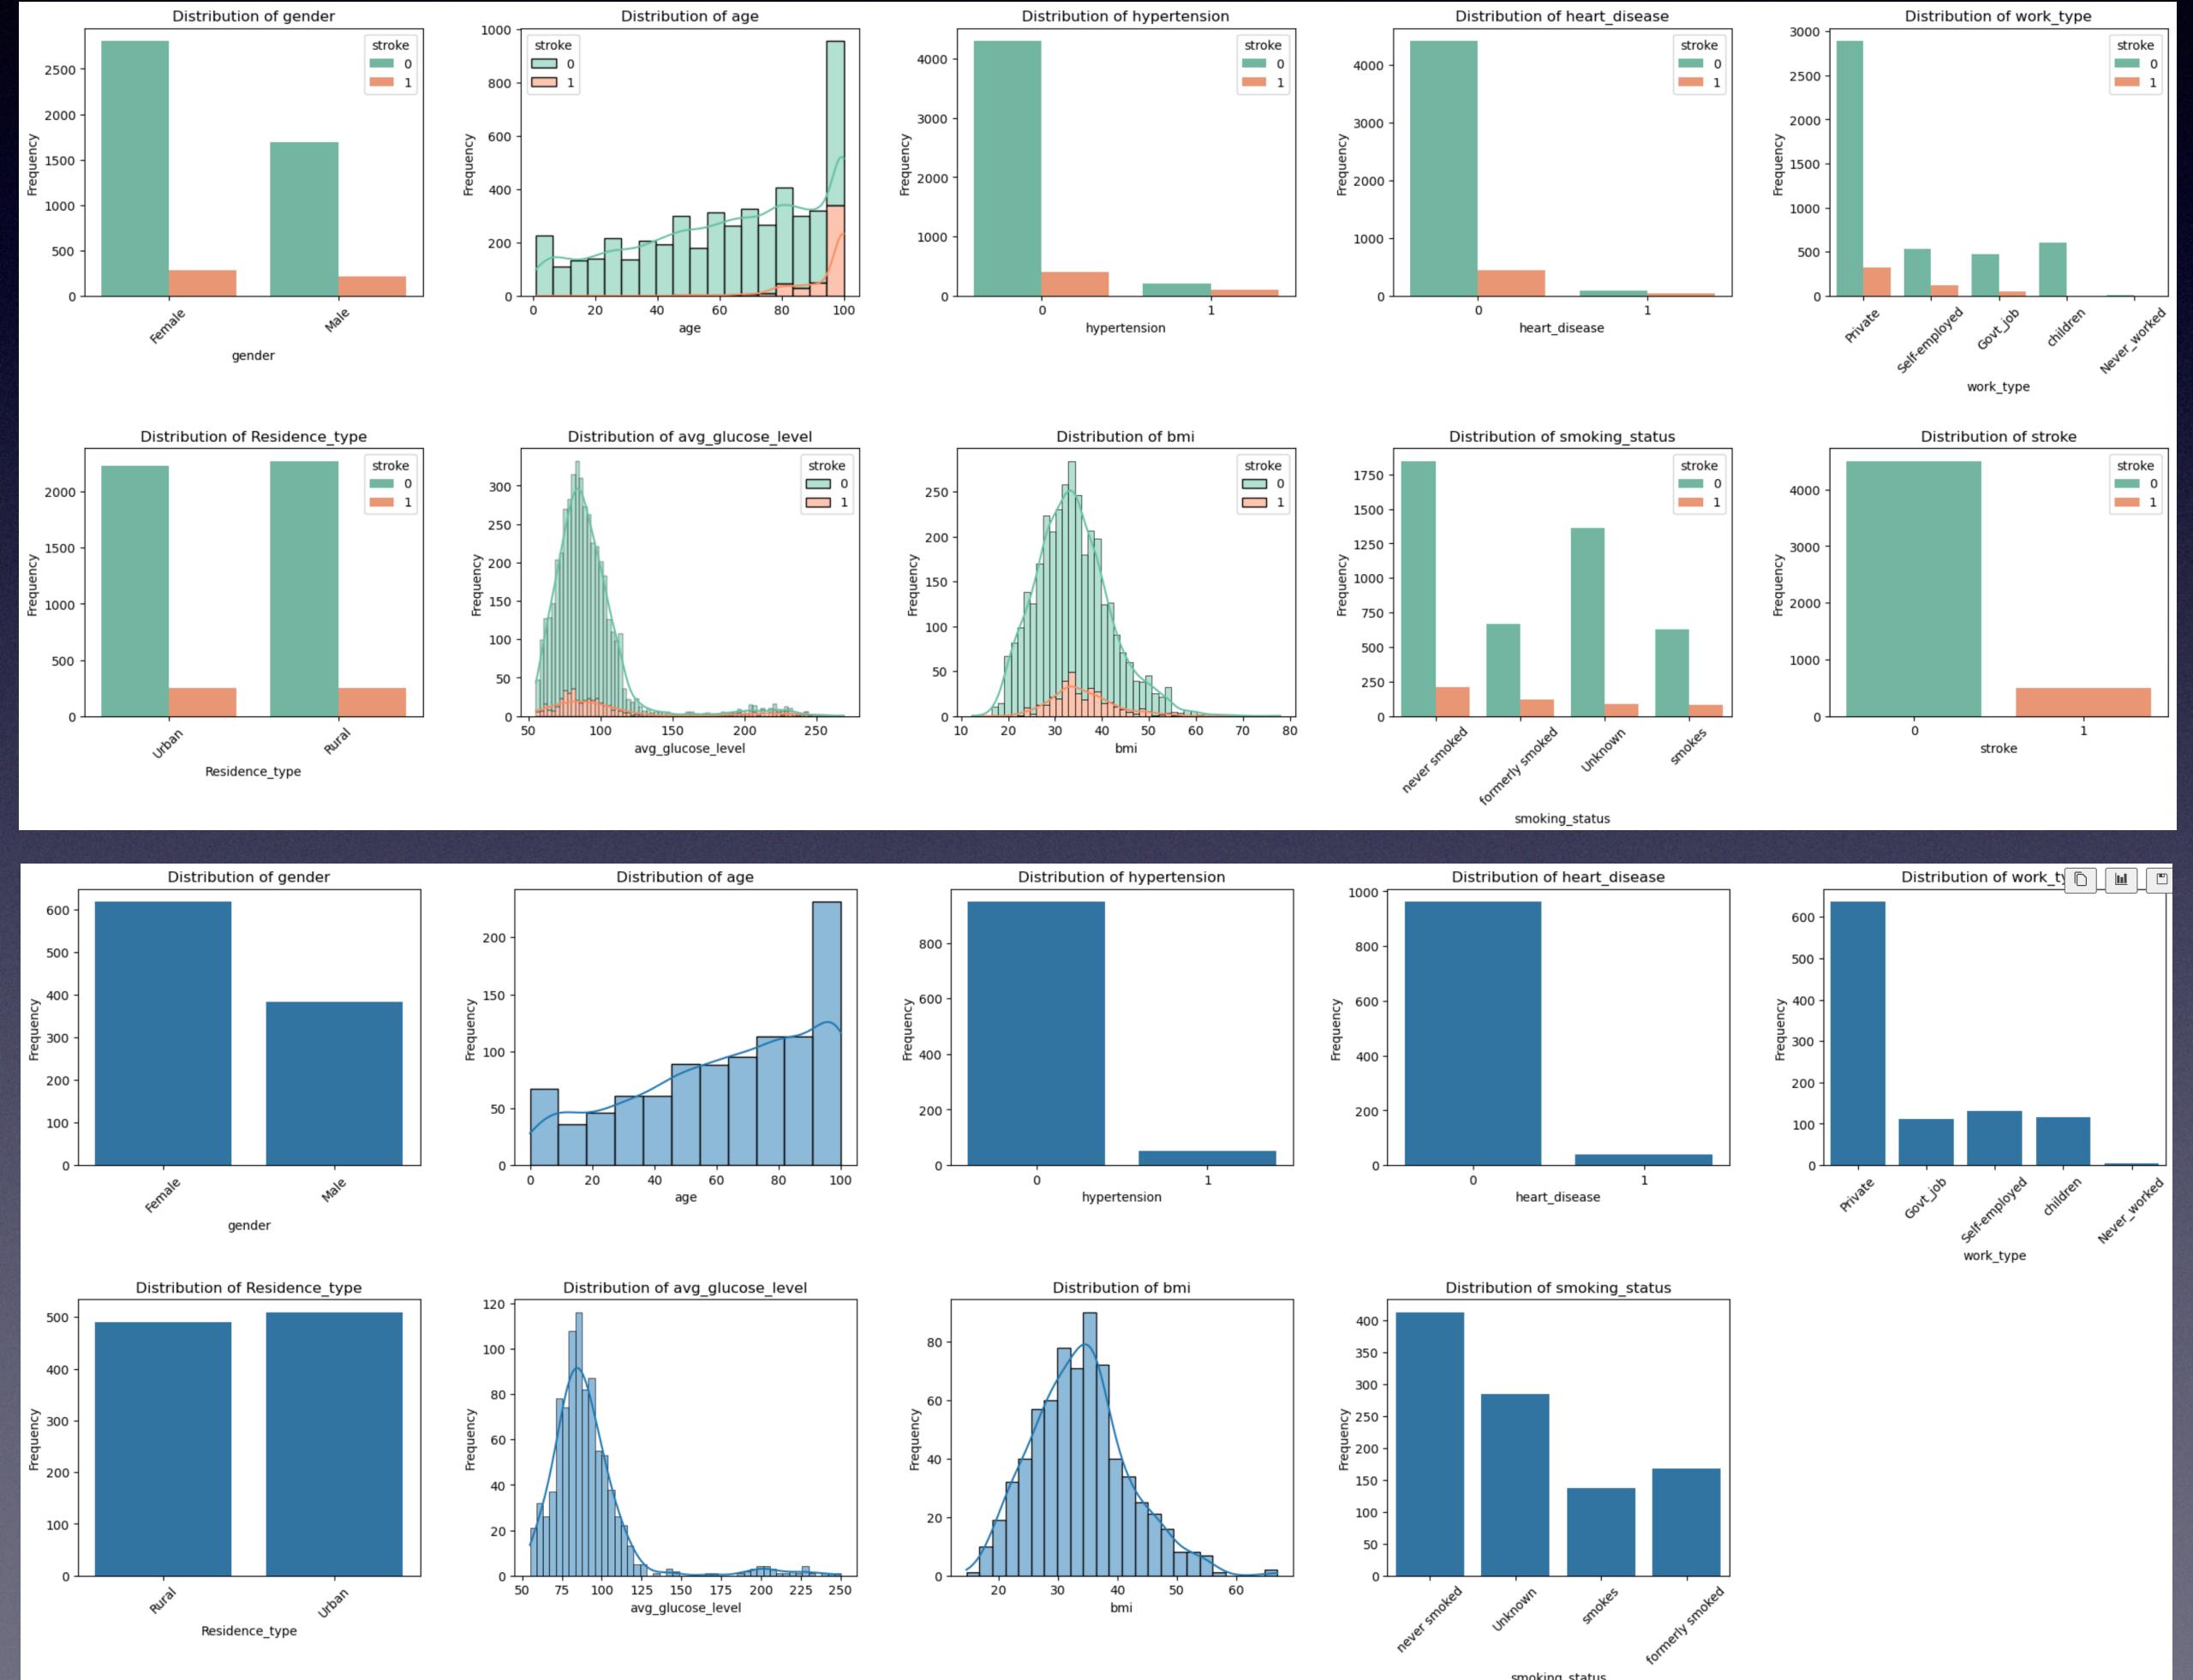
1. Problem Statement

- **Supervised Binary Classification Problem**
- **Predict stroke patients** based on medical and personal data
- **Data Set:**
 - **5000 samples** in train set, **1000 samples** in test set
 - **fictional data**
 - **source:** <https://www.kaggle.com/competitions/ida-ml-1-challenge-summer23/overview>

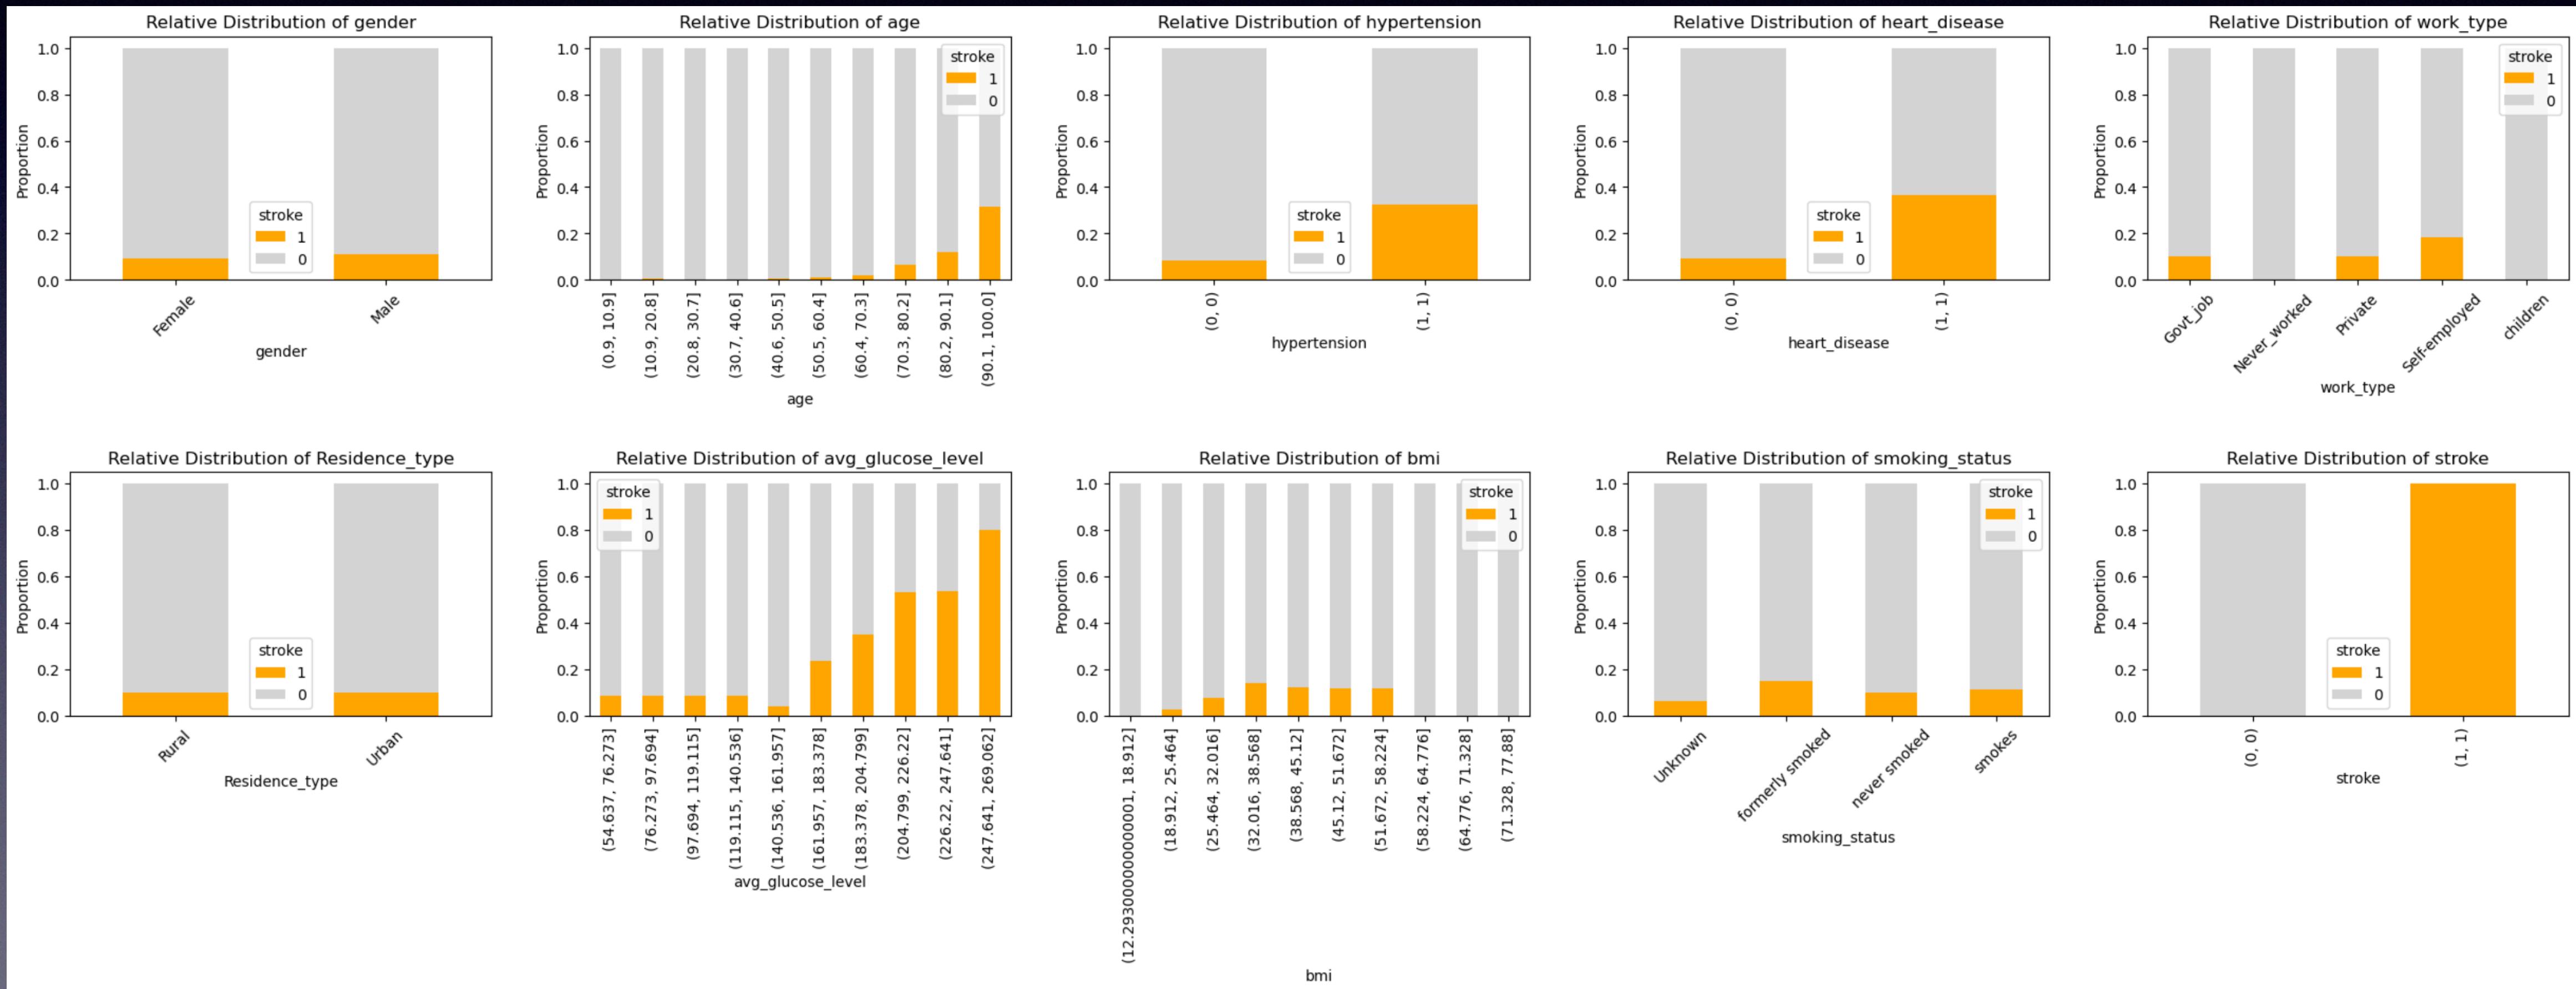
2. Data Analysis & Preprocessing



2. Data Analysis & Preprocessing



2. Data Analysis & Preprocessing



2. Data Analysis & Preprocessing

```
2024-08-04 16:14:33.772 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column gender: 0 / 5000
2024-08-04 16:14:33.772 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column age: 0 / 5000
2024-08-04 16:14:33.773 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column hypertension: 0 / 5000
2024-08-04 16:14:33.774 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column heart_disease: 0 / 5000
2024-08-04 16:14:33.775 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column work_type: 0 / 5000
2024-08-04 16:14:33.775 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column Residence_type: 0 / 5000
2024-08-04 16:14:33.776 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column avg_glucose_level: 393 / 5000
2024-08-04 16:14:33.776 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column bmi: 1552 / 5000
2024-08-04 16:14:33.777 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column smoking_status: 1451 / 5000
2024-08-04 16:14:33.778 | INFO  | src.data_understanding.data_exploration:print_na:133 - Number of NaN values in column stroke: 0 / 5000
2024-08-04 16:14:33.779 | INFO  | src.data_understanding.data_exploration:print_na:138 - Number of data points with NaN values: 2756 / 5000
```

2. Data Analysis & Preprocessing

- **Handling imbalanced data:**
 - Use **weighted cross-entropy loss** as loss function
- **Data preprocessing:**
 - Missing values: **Replace missing values with median**
 - Numerical representation: **Transform discrete features** into binary / one-hot-encoded features
 - **min-max-normalize continuous features**, range [0, 1]

3. Machine Learning Methods

- **Naive Implementation** as baseline model
 - random predictions based on target distribution in training set
(should result in AUC of ~0.5)
 - **Logistic Regression** as linear model
 - train on polynomial features ($d=2$) and principal components, effectively enabling some degree of non-linear separation with regards to the original data
 - **Neural Network** as non-linear model
 - allow model to derive relevant non-linear relations on its own
- ⇒ non-linear separation: polynomials / PCs vs. Neural Network architecture

```
7  class NaiveBaseline:-
8      def __init__(self):
9          super().__init__()
10         self.model = None
11         self.p_true = 0.0
12
13     def fit(
14         self,
15         X_train: npt.ArrayLike,
16         y_train: npt.ArrayLike,
17         X_val: npt.ArrayLike,
18         y_val: npt.ArrayLike,
19     ):-:
20         self.p_true = y_train.mean()
21
22         val_preds = self.predict(X_val)
23         val_accuracy = (val_preds == y_val).mean()
24         print(f"Validation accuracy: {val_accuracy}")
25
26     def predict(self, X: ArrayLike) -> ndarray:
27         """returns a random prediction based on the base probability of the true class, without any training"""
28         return np.random.choice([0, 1], size=len(X), p=[1 - self.p_true, self.p_true])
29
30     def predict_proba(self, X: ArrayLike) -> ndarray:
31         return np.full(len(X), self.p_true)
```

```
7  class NaiveBaseline:-
8      def __init__(self):
9          super().__init__()
10         self.model = None
11         self.p_true = 0.0
12
13     def fit(
14         self,
15         X_train: npt.ArrayLike,
16         y_train: npt.ArrayLike,
17         X_val: npt.ArrayLike,
18         y_val: npt.ArrayLike,
19     ):-:
20         self.p_true = y_train.mean()
21
22         val_preds = self.predict(X_val)
23         val_accuracy = (val_preds == y_val).mean()
24         print(f"Validation accuracy: {val_accuracy}")
25
26     def predict(self, X: ArrayLike) -> ndarray:
27         """returns a random prediction based on the base probability of the true class, without any training"""
28         return np.random.choice([0, 1], size=len(X), p=[1 - self.p_true, self.p_true])
29
30     def predict_proba(self, X: ArrayLike) -> ndarray:
31         return np.full(len(X), self.p_true)
```

```

12 class BinaryLogisticRegression(object):
13     name = "Binary Logistic Regression"
14
15     def __init__(self, n_features: int = 20, epochs: int = 20, learning_rate: float = 0.1, batch_size: int = 8, lambda_reg: float = 0.1, regularization: str = "L2"):
16         self.n_features = n_features
17         self.epochs = epochs
18         self.learning_rate = learning_rate
19         self.batch_size = batch_size
20         self.lambda_reg = lambda_reg
21         self.regularization = regularization
22
23     X: (n_samples, n_features)
24     Y: (n_samples,)
25
26     self.W: (n_features, 1)
27     self.B: (1,)
28
29     self.class_weights = np.zeros(2)
30
31     # initialize weights and bias to zeros
32     self.W = np.zeros(n_features)
33     self.B = 0.0
34
35     def forward(self, X: npt.ArrayLike) -> npt.ArrayLike:
36         Compute forward pass of the logistic regression model.
37
38         Y_hat: (n_samples,)
39
40         z = np.dot(X, self.W) + self.B
41         Y_hat = sigmoid(z)
42         return Y_hat
43
44     def predict(self, X: npt.ArrayLike) -> npt.ArrayLike:
45         Create a prediction matrix with `self.forward()`
46
47         pred: (n_samples,)
48
49         y_hat = self.forward(X)
50         pred = np.where(y_hat >= 0.5, 1, 0)
51
52         return pred
53
54     def predict_proba(self, X: npt.ArrayLike) -> npt.ArrayLike:
55         Predict probabilities for the input data.
56
57         Y_hat: (n_samples,)
58
59         Y_hat = self.forward(X)
60         return Y_hat
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160

```

```

class BinaryLogisticRegression(object):
    name = "Binary Logistic Regression"

    def __init__(self, n_features: int = 20, epochs: int = 20, learning_rate: float = 0.1, batch_size: int = 8, lambda_reg: float = 0.1, regularization: str = "L2"):
        """
        X: (n_samples, n_features)
        Y: (n_samples,)

        self.W: (n_features, 1)
        self.B: (1,)
        """
        self.n_features = n_features
        self.epochs = epochs
        self.learning_rate = learning_rate
        self.batch_size = batch_size
        self.lambda_reg = lambda_reg
        self.regularization = regularization

        self.class_weights = np.zeros(2)

        # initialize weights and bias to zeros
        self.W = np.zeros(n_features)
        self.B = 0.0

    def forward(self, X: npt.ArrayLike) -> npt.ArrayLike:
        """
        Compute forward pass of the logistic regression model.

        Y_hat: (n_samples,)
        """
        z = np.dot(X, self.W) + self.B
        Y_hat = sigmoid(z)
        return Y_hat

    def predict(self, X: npt.ArrayLike) -> npt.ArrayLike:
        """
        Create a prediction matrix with `self.forward()`

        pred: (n_samples,)
        """
        y_hat = self.forward(X)
        pred = np.where(y_hat >= 0.5, 1, 0)

        return pred

    def predict_proba(self, X: npt.ArrayLike) -> npt.ArrayLike:
        """
        Predict probabilities for the input data.

        Y_hat: (n_samples,)
        """
        Y_hat = self.forward(X)
        return Y_hat

    def _backward(self, X: npt.ArrayLike, Y: npt.ArrayLike, Y_hat: npt.ArrayLike) -> Tuple[npt.ArrayLike, npt.ArrayLike]:
        """
        Compute back propagation for logistic regression.

        batch_size = X.shape[0]

        # derivative of the loss with respect to the pre-activation of the output layer
        # equivalent to d L d z = Y_hat - Y in non-weighted binary cross-entropy
        d_L_d_z = delta_weighted_bce(Y_hat, Y, self.class_weights)

        d_L_d_W = np.dot(X.T, d_L_d_z) / batch_size
        d_L_d_B = np.sum(d_L_d_z, axis=0) / batch_size

        if self.regularization == "L1":
            # add L1 regularization derivative
            d_L_d_W += self.lambda_reg * np.sign(self.W)
        else:
            # add L2 regularization derivative
            d_L_d_W += self.lambda_reg * self.W
        """
        return d_L_d_W, d_L_d_B

    def fit(self, X: npt.ArrayLike, Y: npt.ArrayLike, X_val: npt.ArrayLike, Y_val: npt.ArrayLike, plot: bool = False):
        """
        Fit the logistic regression model to the training data.

        n = X.shape[0]

        self.class_weights = compute_class_weight(class_weight="balanced", classes=np.array([0, 1]), y=Y)
        """

        losses_train = []
        losses_val = []

        for epoch in range(self.epochs):
            epoch_loss_train = 0.0
            for i in range(0, n, self.batch_size):
                X_batch = X[i : i + self.batch_size]
                Y_batch = Y[i : i + self.batch_size]

                Y_hat = self.forward(X_batch)
                d_L_d_W, d_L_d_B = self._backward(X_batch, Y_batch, Y_hat)

                self.W -= self.learning_rate * d_L_d_W
                self.B -= self.learning_rate * d_L_d_B

                batch_loss = weighted_binary_cross_entropy_loss(Y_hat, Y_batch, self.class_weights)
                epoch_loss_train += batch_loss * len(Y_batch)

            loss_train = epoch_loss_train / n

            Y_hat_val = self.forward(X_val)
            loss_val = weighted_binary_cross_entropy_loss(Y_hat_val, Y_val, self.class_weights)

            losses_train.append(loss_train)
            losses_val.append(loss_val)

            if plot:
                print(f"Epoch {epoch}: Train Loss: {loss_train}, Val Loss: {loss_val}")

        if plot:
            plt.plot(losses_train, label="Training Loss")
            plt.plot(losses_val, label="Validation Loss")
            plt.xlabel("Number of epochs")
            plt.ylabel("Loss")
            plt.title("Training and Validation Loss")
            plt.legend()
            plt.show()
            print("Training complete.")

```

```

1 You, 1 hour ago | 1 author (You)
2 class BinaryNeuralNetwork(object):
3     name = "Binary Neural Network"
4
5     def __init__(self,
6         n_features: int,
7         n_hidden_units: int,
8         n_hidden_layers: int,
9         epochs: int = 10,
10        learning_rate: float = 0.05,
11        batch_size: int = 32,
12        lambda_reg: float = 0.01,
13    ):
14        """
15            X: n_input x d_input
16            Y: n_input
17
18            self.W[0]: n_features x n_hidden_units
19            self.B[0]: n_hidden_units
20
21            self.W[i: 0 < i < (n_hidden-1)]: n_hidden_units x n_hidden_units
22            self.B[i: 0 < i < (n_hidden-1)]: n_hidden_units
23
24            self.W[-1]: n_hidden_units
25            self.B[-1]: 1
26
27            self.n_features = n_features
28            self.n_hidden_units = n_hidden_units
29            self.n_hidden_layers = n_hidden_layers
30            self.epochs = epochs
31            self.learning_rate = learning_rate
32            self.batch_size = batch_size
33            self.lambda_reg = lambda_reg
34
35            self.class_weights = np.zeros(2)
36
37            self.W = []
38            self.B = []
39
40            self.initialize_params()
41
42        def initialize_params(self):
43            """
44                Initialize weights and biases (He initialization)
45            for i in range(self.n_hidden_layers + 1):
46                if i == 0: # input -> h_1
47                    limit = np.sqrt(2 / (self.n_features))
48                    w = np.random.normal(0, limit, (self.n_features, self.n_hidden_units))
49                    b = np.zeros((self.n_hidden_units,))
50
51                elif i < self.n_hidden_layers: # h_i -> h_{i+1}
52                    limit = np.sqrt(2 / (self.n_hidden_units))
53                    w = np.random.normal(
54                        0, limit, (self.n_hidden_units, self.n_hidden_units)
55                    )
56                    b = np.zeros((self.n_hidden_units,))
57
58                else: # h_{n} -> output
59                    limit = np.sqrt(2 / (self.n_hidden_units))
60                    w = np.random.normal(0, limit, (self.n_hidden_units, 1))
61                    b = np.zeros((1,))
62
63                    self.W.append(w)
64                    self.B.append(b)
65
66
67
68
69
70
71
72

```

```

12 class BinaryNeuralNetwork(object):
13
14     def forward(self, X: npt.ArrayLike, return_intermediates=False) -> npt.ArrayLike:
15         """
16             Compute forward pass of the neural network.
17
18             Y_hat: n_samples
19
20             A = [X]
21             Z = []
22
23             # forward propagation
24             a_i = X
25             for i in range(self.n_hidden_layers + 1):
26                 z_i = np.dot(a_i, self.W[i]) + self.B[i] # add bias row-wise
27
28                 if i < self.n_hidden_layers:
29                     a_i = relu(z_i)
30                 else: # output layer
31                     z_i = z_i.squeeze() # transform (n_samples, 1) to (n_samples,)
32                     a_i = sigmoid(z_i) # (n_samples,)
33
34                     Z.append(z_i)
35                     A.append(a_i)
36
37             Y_hat = A[-1].squeeze()
38
39             if return_intermediates:
40                 return Y_hat, A, Z
41
42             return Y_hat
43
44     def predict(self, X: npt.ArrayLike) -> npt.ArrayLike:
45         """
46             Create a prediction matrix with `self.forward()`
47
48             pred: (n_samples,)
49
50             y_hat = self.forward(X)
51             pred = np.where(y_hat >= 0.5, 1, 0)
52
53             return pred
54
55     def predict_proba(self, X: npt.ArrayLike) -> npt.ArrayLike:
56         """
57             Create a prediction matrix with `self.forward()`
58
59             y_hat: (n_samples,)
60
61             y_hat = self.forward(X)
62
63             return y_hat
64
65
66
67
68
69
70
71
72

```

```

1 You, 1 hour ago | 1 author (You)
2 class BinaryNeuralNetwork(object):
3     name = "Binary Neural Network"
4
5     def __init__(self,
6         n_features: int,
7         n_hidden_units: int,
8         n_hidden_layers: int,
9         epochs: int = 10,
10        learning_rate: float = 0.05,
11        batch_size: int = 32,
12        lambda_reg: float = 0.01,
13    ):
14        """
15            X: n_input x d_input
16            Y: n_input
17
18            self.W[0]: n_features x n_hidden_units
19            self.B[0]: n_hidden_units
20
21            self.W[i: 0 < i < (n_hidden-1)]: n_hidden_units x n_hidden_units
22            self.B[i: 0 < i < (n_hidden-1)]: n_hidden_units
23
24            self.W[-1]: n_hidden_units
25            self.B[-1]: 1
26
27            self.n_features = n_features
28            self.n_hidden_units = n_hidden_units
29            self.n_hidden_layers = n_hidden_layers
30            self.epochs = epochs
31            self.learning_rate = learning_rate
32            self.batch_size = batch_size
33            self.lambda_reg = lambda_reg
34
35            self.class_weights = np.zeros(2)
36
37            self.W = []
38            self.B = []
39
40            self.initialize_params()
41
42        def initialize_params(self):
43            """
44                Initialize weights and biases (He initialization)
45            for i in range(self.n_hidden_layers + 1):
46                if i == 0: # input -> h_1
47                    limit = np.sqrt(2 / (self.n_features))
48                    w = np.random.normal(0, limit, (self.n_features, self.n_hidden_units))
49                    b = np.zeros((self.n_hidden_units,))
50
51                elif i < self.n_hidden_layers: # h_i -> h_{i+1}
52                    limit = np.sqrt(2 / (self.n_hidden_units))
53                    w = np.random.normal(
54                        0, limit, (self.n_hidden_units, self.n_hidden_units)
55                    )
56                    b = np.zeros((self.n_hidden_units,))
57
58                else: # h_{n} -> output
59                    limit = np.sqrt(2 / (self.n_hidden_units))
60                    w = np.random.normal(0, limit, (self.n_hidden_units, 1))
61                    b = np.zeros((1,))
62
63                self.W.append(w)
64                self.B.append(b)
65
66
67
68
69
70
71
72

```

```

12 class BinaryNeuralNetwork(object):
13
14     def forward(self, X: npt.ArrayLike, return_intermediates=False) -> npt.ArrayLike:
15         """
16             Compute forward pass of the neural network.
17
18             Y_hat: n_samples
19
20             A = [X]
21             Z = []
22
23             # forward propagation
24             a_i = X
25             for i in range(self.n_hidden_layers + 1):
26                 z_i = np.dot(a_i, self.W[i]) + self.B[i] # add bias row-wise
27
28                 if i < self.n_hidden_layers:
29                     a_i = relu(z_i)
30                 else: # output layer
31                     z_i = z_i.squeeze() # transform (n_samples, 1) to (n_samples,)
32                     a_i = sigmoid(z_i) # (n_samples,)
33
34             Z.append(z_i)
35             A.append(a_i)
36
37             Y_hat = A[-1].squeeze()
38
39             if return_intermediates:
40                 return Y_hat, A, Z
41
42             return Y_hat
43
44     def predict(self, X: npt.ArrayLike) -> npt.ArrayLike:
45         """
46             Create a prediction matrix with `self.forward()`
47
48             pred: (n_samples,)
49
50             y_hat = self.forward(X)
51             pred = np.where(y_hat >= 0.5, 1, 0)
52
53             return pred
54
55     def predict_proba(self, X: npt.ArrayLike) -> npt.ArrayLike:
56         """
57             Create a prediction matrix with `self.forward()`
58
59             y_hat: (n_samples,)
60
61             y_hat = self.forward(X)
62
63             return y_hat
64
65
66
67
68
69
70
71
72

```

```

12 class BinaryNeuralNetwork(object):-
13
14     def _backward(-
15         self,-
16         Y: npt.ArrayLike,-
17         Y_hat: npt.ArrayLike,-
18         A: npt.ArrayLike,-
19         Z: npt.ArrayLike,-
20     ) -> Tuple[npt.ArrayLike, npt.ArrayLike]:-
21         """
22             Compute back propagation of the neural network.
23         """
24         batch_size = Y.shape[0]
25
26         d_L_d_W = [np.zeros_like(w) for w in self.W]
27         d_L_d_B = [np.zeros_like(b) for b in self.B]
28
29         # derivative of the loss with respect to the pre-activation of the output layer
30         d_L_d_z = delta_weighted_bce(Y_hat, Y, self.class_weights)
31         d_L_d_z = d_L_d_z.reshape(-1, 1) # reshape from (n_samples,) to (n_samples, 1)
32
33         # derivatives of the loss with respect to the weights and biases of the last layer
34         d_L_d_W[-1] = np.dot(A[-2].T, d_L_d_z) / batch_size
35         d_L_d_B[-1] = np.sum(d_L_d_z, axis=0) / batch_size
36
37         # back propagation using chain rule
38         for l in range(self.n_hidden_layers - 1, -1, -1):
39             d_L_d_z = np.dot(d_L_d_z, self.W[l + 1].T) * relu_prime(
40                 Z[l]
41             ) #  $d^l = (d^{l+1})^T \times W^{l+1} \times (dA^l / dz^l)$ 
42
43             d_L_d_W[l] = (
44                 np.dot(A[l].T, d_L_d_z) / batch_size
45             ) # note: A[l] is actually correct here - A[0] is X!
46             d_L_d_B[l] = np.sum(d_L_d_z, axis=0) / batch_size
47
48             # add L2 regularization derivative
49             d_L_d_W[l] += self.lambda_reg * self.W[l]
50
51         return d_L_d_W, d_L_d_B
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221

```

```

12 class BinaryNeuralNetwork(object):-
13
14     def _backward(-
15         self,-
16         Y: npt.ArrayLike,-
17         Y_hat: npt.ArrayLike,-
18         A: npt.ArrayLike,-
19         Z: npt.ArrayLike,-
20     ) -> Tuple[npt.ArrayLike, npt.ArrayLike]:-
21         """
22             Compute back propagation of the neural network.
23         """
24         batch_size = Y.shape[0]
25
26         d_L_d_W = [np.zeros_like(w) for w in self.W]
27         d_L_d_B = [np.zeros_like(b) for b in self.B]
28
29         # derivative of the loss with respect to the pre-activation of the output layer
30         d_L_d_z = delta_weighted_bce(Y_hat, Y, self.class_weights)
31         d_L_d_z = d_L_d_z.reshape(-1, 1) # reshape from (n_samples,) to (n_samples, 1)
32
33         # derivatives of the loss with respect to the weights and biases of the last layer
34         d_L_d_W[-1] = np.dot(A[-2].T, d_L_d_z) / batch_size
35         d_L_d_B[-1] = np.sum(d_L_d_z, axis=0) / batch_size
36
37
38         # back propagation using chain rule
39         for l in range(self.n_hidden_layers - 1, -1, -1):
40             d_L_d_z = np.dot(d_L_d_z, self.W[l + 1].T) * relu_prime(
41                 Z[l]
42             ) #  $d^l = (d^{l+1})^T \times W^{l+1} \times (dA^l / dz^l)$ 
43
44             d_L_d_W[l] = (
45                 np.dot(A[l].T, d_L_d_z) / batch_size
46             ) # note: A[l] is actually correct here - A[0] is X!
47             d_L_d_B[l] = np.sum(d_L_d_z, axis=0) / batch_size
48
49             # add L2 regularization derivative
50             d_L_d_W[l] += self.lambda_reg * self.W[l]
51
52
53         return d_L_d_W, d_L_d_B
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221

```

```
5  def weighted_binary_cross_entropy_loss(←
6  |     output, target, class_weights=None, epsilon=0.001←
7  ):←
8  """←
9      Compute the weighted binary cross-entropy loss.←
10     """←
11     if class_weights is None:←
12         class_weights = compute_class_weights(target)←
13     ←
14     # avoiding numerical instability←
15     output = np.clip(output, epsilon, 1 - epsilon)←
16     ←
17     loss = class_weights[1] * (target * np.log(output)) + class_weights[0] * (←
18     |     (1 - target) * np.log(1 - output)←
19     )←
20     ←
21     return -np.mean(loss)←
22     ←
23     ←
24     def compute_class_weights(y: npt.ArrayLike) -> npt.ArrayLike:←
25     """←
26         Compute class weights for imbalanced datasets.←
27     """←
28     y = np.array(y, dtype=np.int64)←
29     classes = np.unique(y)←
30     class_counts = np.bincount(y)←
31     total_samples = len(y)←
32     ←
33     class_weights = {←
34         cls: total_samples / (len(classes) * count)←
35         for cls, count in zip(classes, class_counts)←
36     }←
37     ←
38     weights_array = np.array([class_weights[cls] for cls in classes], dtype=np.float64)←
39     ←
40     return weights_array
```

```
5  def weighted_binary_cross_entropy_loss(←
6  |     output, target, class_weights=None, epsilon=0.001←
7  ):←
8  """←
9      Compute the weighted binary cross-entropy loss.←
10     """←
11     if class_weights is None:←
12         class_weights = compute_class_weights(target)←
13     ←
14     # avoiding numerical instability←
15     output = np.clip(output, epsilon, 1 - epsilon)←
16     ←
17     loss = class_weights[1] * (target * np.log(output)) + class_weights[0] * (←
18     |     (1 - target) * np.log(1 - output)←
19     )←
20     ←
21     return -np.mean(loss)←
22     ←
23     ←
24     def compute_class_weights(y: npt.ArrayLike) -> npt.ArrayLike:←
25     """←
26         Compute class weights for imbalanced datasets.←
27     """←
28     y = np.array(y, dtype=np.int64)←
29     classes = np.unique(y)←
30     class_counts = np.bincount(y)←
31     total_samples = len(y)←
32     ←
33     class_weights = {←
34         cls: total_samples / (len(classes) * count)←
35         for cls, count in zip(classes, class_counts)←
36     }←
37     ←
38     weights_array = np.array([class_weights[cls] for cls in classes], dtype=np.float64)←
39     ←
40     return weights_array
```

4. Feature Selection for Logistic Regression

- Polynomial Features, Primal Decision Function
- Principal Components through Principal Component Analysis
- Feature Selection using custom backward selection protocol

4. Feature Selection for Logistic Regression

- Feature Selection Protocol: Backward Selection
 - delete one feature per iteration
 - train LogisticRegression model, calculate risk estimate on validation set at each step
 - as soon as all but one features are deleted: return the features that the model with the lowest risk estimate was trained on

```
42 def feature_selection(-
43     X,
44     y,
45     n_features_per_iteration: int = 1,
46     plot: bool = False,
47 ):-:
48     n_features = X.shape[1]
49     deleted_features = []
50     optimal_number_of_features_to_delete = 0
51
52     risk_estimates = [np.inf]
53     lowest_risk_estimate = np.inf
54
55     X_tmp = X.copy()
56     y_tmp = y.copy()
57
58     original_indices = np.arange(n_features)
59
60     for i in range(0, n_features - 1):
61         X_train, X_val, y_train, y_val = train_test_split(
62             X_tmp, y_tmp, test_size=0.2, random_state=42
63         )
64         model = LogisticRegression(
65             penalty="l1",
66             solver="liblinear",
67             max_iter=1000,
68             class_weight="balanced",
69         )
69         log_reg_tmp = model.fit(X_train, y_train)
70         y_pred = log_reg_tmp.predict_proba(X_val)[:, 1]
71         risk_estimate = weighted_binary_cross_entropy_loss(y_pred, y_val)
72         log_message = f"i = {i} / {n_features-1} - Empirical Risk Estimate: {risk_estimate}, Previous Estimate: {risk_estimates[-1]}"
73
74         features_to_delete = np.argsort(np.abs(log_reg_tmp.coef_))[0][
75             :n_features_per_iteration
76         ]
77         deleted_features.extend(original_indices[features_to_delete].tolist())
78
79         X_tmp = np.delete(X_tmp, features_to_delete, axis=1)
80         original_indices = np.delete(original_indices, features_to_delete)
81
82         risk_estimates.append(risk_estimate)
83
84         if risk_estimate < lowest_risk_estimate:
85             lowest_risk_estimate = risk_estimate
86             optimal_number_of_features_to_delete = i
87             log_message += " - new lowest risk estimate"
88
89         print(log_message)
90
91     features_to_delete = deleted_features[:optimal_number_of_features_to_delete]
92
93     if plot:
94         print(
95             f"Number of features to delete: {optimal_number_of_features_to_delete} / {n_features}"
96         )
97
98         plt.plot(
99             range(
100                 0, n_features - 1 + n_features_per_iteration, n_features_per_iteration
101             ),
102             risk_estimates,
103             label="Training Loss",
104         )
105         plt.xlabel("Number of deleted features")
106         plt.ylabel("Weighted Binary Cross Entropy Loss")
107         plt.title("Risk estimate as a function of deleted features")
108         plt.legend()
109         plt.show()
110
111
112     return features_to_delete, risk_estimates
```

5. Model Selection & Evaluation Protocol

- Nested Cross-Validation protocol for hyper parameter tuning and evaluation
 - 5-fold outer loop, 5-fold inner loop
 - inner loop: iterate over all given hyper parameter sets, calculate empirical risk estimate per iteration using 5-fold cross-validation
 - outer loop: 5-fold split of the data
 - after outer loop: average empirical risk estimates per hyper parameter set, pick set with lowest empirical risk estimate, evaluate model using 5-fold cross-validation, fit final model on whole training set

```
1 param_grid_lr = {  
2     "n_features": [X_selected.shape[1]],  
3     "epochs": [10, 20, 30],  
4     "learning_rate": [0.1, 0.05, 0.01],  
5     "lambda_reg": [0.1, 0.01, 0.001],  
6 }
```

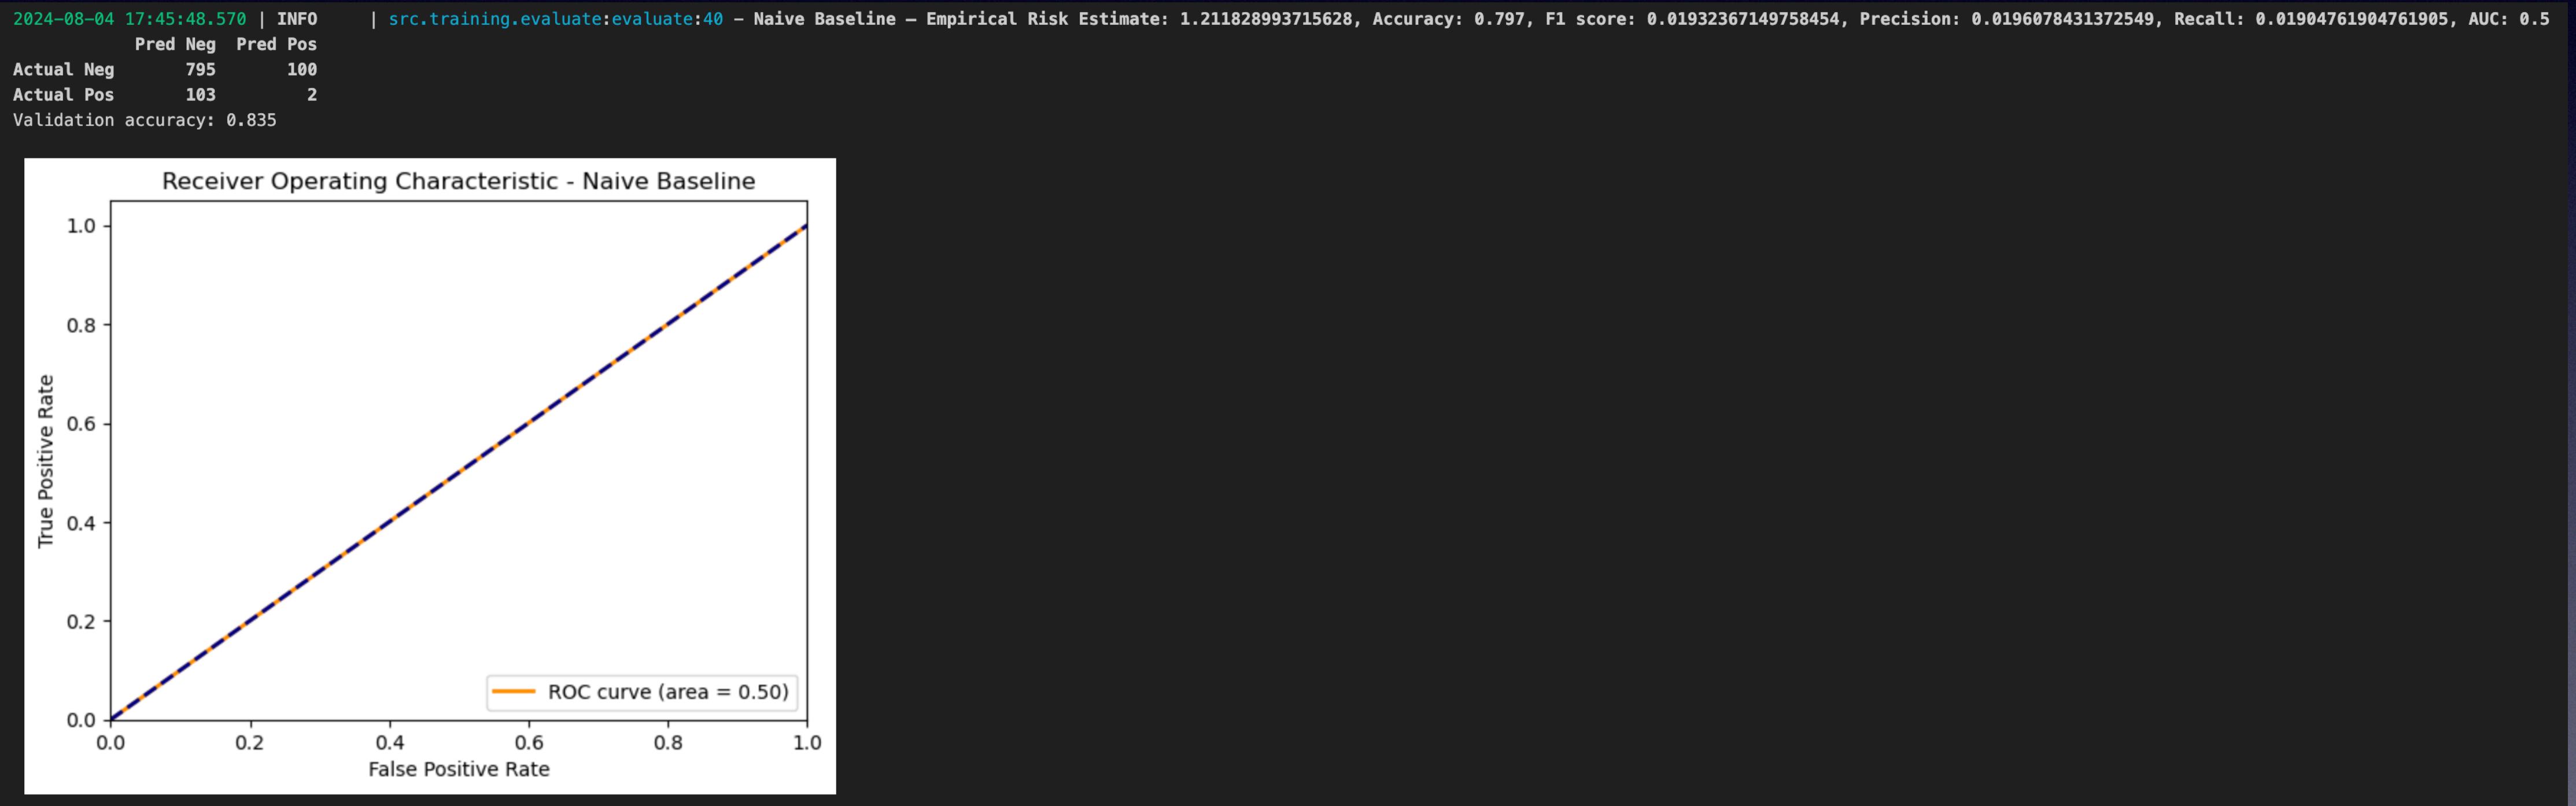
```
1 param_grid_nn = {  
2     "n_features": [X_original.shape[1]],  
3     "epochs": [10, 20, 30],  
4     "learning_rate": [0.1, 0.05, 0.01],  
5     "lambda_reg": [0.1, 0.01, 0.001],  
6     "n_hidden_units": [16, 32, 64],  
7     "n_hidden_layers": [1, 2],  
8 }
```

```

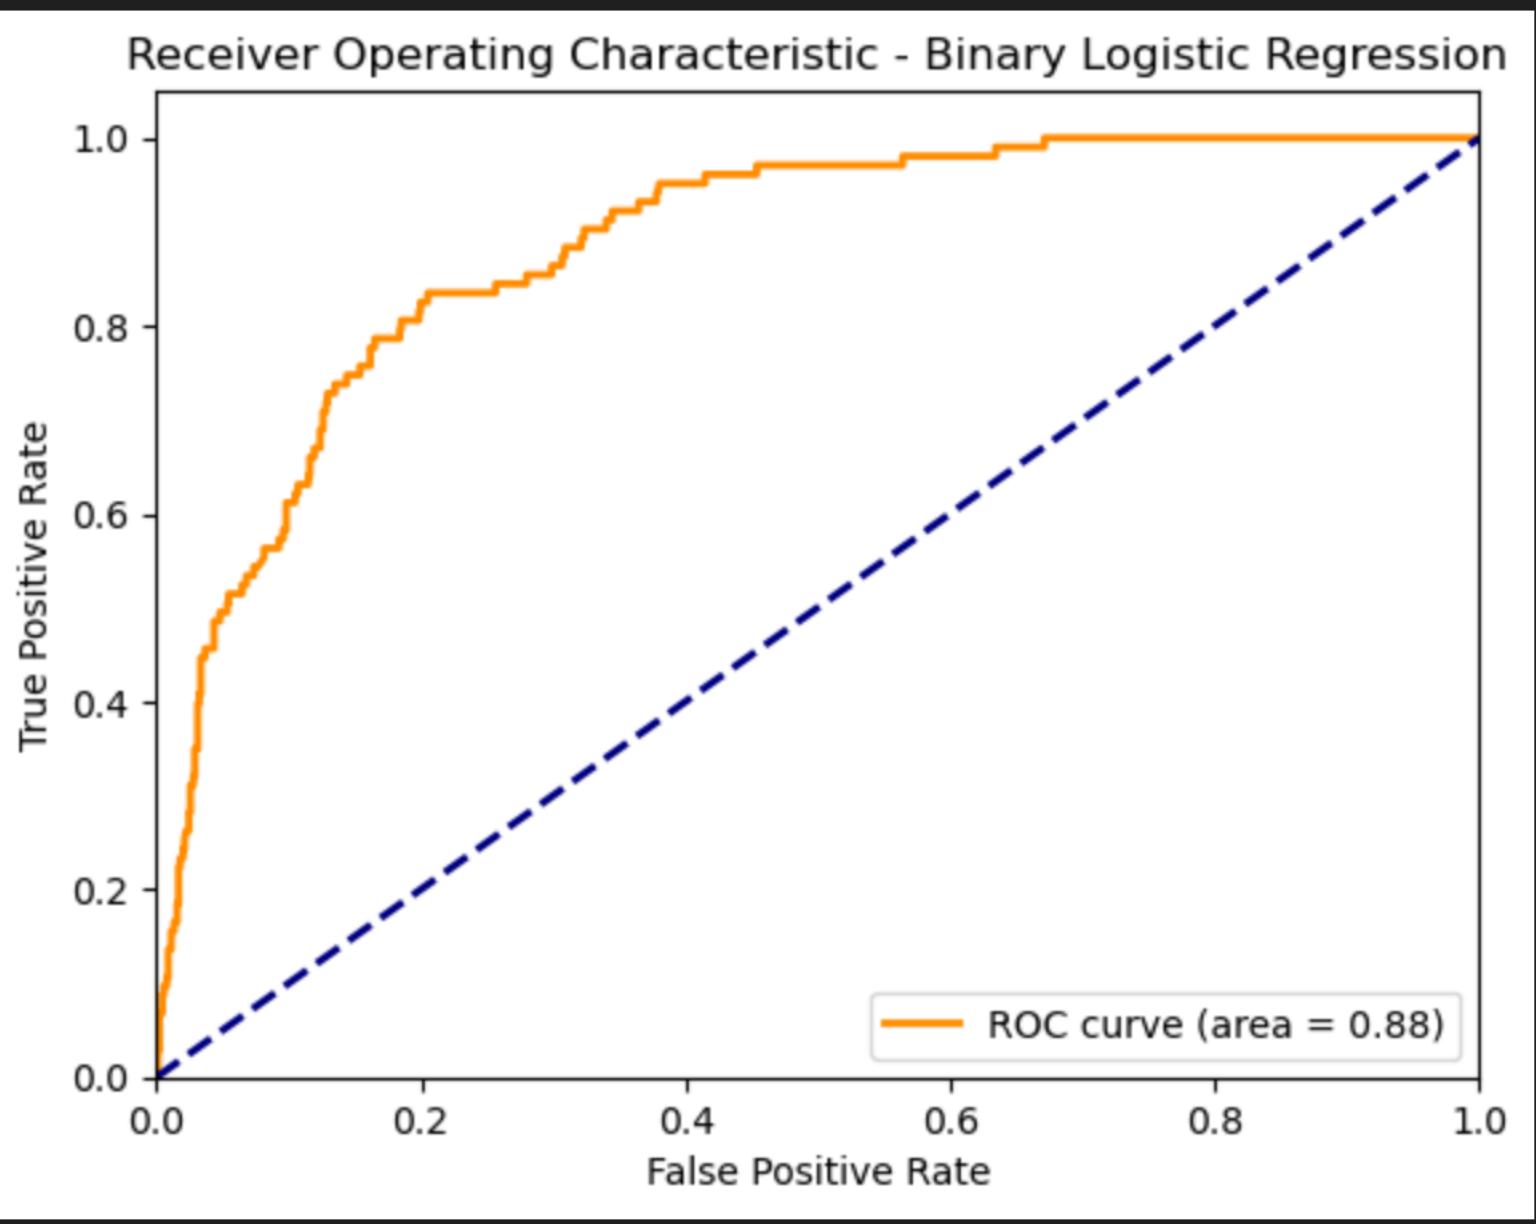
146 def nested_cross_validation(X, y, Model, param_grid, k=5):
147     print(
148         f"Performing nested cross-validation for model {Model.name} with {X.shape[0]} samples"
149     )
150     # list of one dict per parameter combination
151     param_combinations = [
152         dict(zip(param_grid.keys(), values)) for values in product(*param_grid.values())
153     ]
154     n_combinations = len(param_combinations)
155
156     R_est = np.zeros((k, len(param_combinations)))
157
158     kf_outer = KFold(n_splits=k, shuffle=True)
159
160     for i, (outer_train_idx, outer_test_idx) in enumerate(kf_outer.split(X)):
161         X_train_outer, X_val_outer = X[outer_train_idx], X[outer_test_idx]
162         y_train_outer, y_val_outer = y[outer_train_idx], y[outer_test_idx]
163
164         for j, params in enumerate(param_combinations):
165             model = Model(**params)
166
167             R_est[i, j] = k_fold_cross_validation(
168                 model, X_train_outer, y_train_outer, k=k, fit_final_model=False
169             )
170             print(
171                 f"Empirical risk estimate for fold {i+1} / {k}, parameter set {j+1} / {n_combinations}: {R_est[i, j]}",
172                 end="\r",
173             )
174
175     R_est_params = np.mean(R_est, axis=0)
176
177     params = param_combinations[np.argmin(R_est_params)]
178     print(
179         f"Risk estimate for argmin(R_est_params): {R_est_params[np.argmin(R_est_params)]}"
180     )
181     print(f"Selected best parameters: {params}")
182
183     model_final = Model(**params)
184     model_final_trained, R_est, acc, f1, auc, prec, rec, cm_df, fpr, tpr = (
185         k_fold_cross_validation(model_final, X, y, k=k, fit_final_model=True)
186     )
187     print(
188         f"{Model.name} - Empirical Risk Estimate: {R_est}, Accuracy: {acc}, F1 score: {f1}, Precision: {prec}, Recall: {rec}, AUC: {auc}\n{cm_df}"
189     )
190
191     # Plot ROC curve
192     plt.figure()
193     plt.plot(fpr, tpr, color="darkorange", lw=2, label=f"ROC curve (area = {auc:0.2f})")
194     plt.plot([0, 1], [0, 1], color="navy", lw=2, linestyle="--")
195     plt.xlim([0.0, 1.0])
196     plt.ylim([0.0, 1.05])
197     plt.xlabel("False Positive Rate")
198     plt.ylabel("True Positive Rate")
199     plt.title(f"Receiver Operating Characteristic - {Model.name}")
200     plt.legend(loc="lower right")
201     plt.show()
202
203     return model_final_trained, params, R_est, acc, f1, auc

```

6. Results – Naive Baseline



6. Results – Logistic Regression

```
[49] ✓ 2m 5.4s
...
... Performing nested cross-validation for model Binary Logistic Regression with 5000 samples
Risk estimate for argmin(R_ests_params): 0.431203867143375447: 0.44956165679782895
Selected best parameters: {'n_features': 3, 'epochs': 20, 'learning_rate': 0.1, 'lambda_reg': 0.001}
Binary Logistic Regression - Empirical Risk Estimate: 0.4326625244737709, Accuracy: 0.7598, F1 score: 0.4106326848282034, Precision: 0.27178922773035485, Recall: 0.8410820460604125, AUC: 0.8783259164403588
      Pred Neg  Pred Pos
Actual Neg      678      219
Actual Pos       17       86
...
...
Receiver Operating Characteristic - Binary Logistic Regression


The figure is a Receiver Operating Characteristic (ROC) plot titled "Receiver Operating Characteristic - Binary Logistic Regression". The x-axis is labeled "False Positive Rate" and ranges from 0.0 to 1.0. The y-axis is labeled "True Positive Rate" and ranges from 0.0 to 1.0. A solid orange line represents the "ROC curve (area = 0.88)". A dashed blue diagonal line from (0,0) to (1,1) represents a random classifier. The ROC curve is above the diagonal line, indicating better performance.


```

6. Results – Neural Network

```
[50] ✓ 20m 28.7s
...
... Performing nested cross-validation for model Binary Neural Network with 5000 samples
Risk estimate for argmin(R_est_params): 0.43757705619196885162: 0.45009287877300624
Selected best parameters: {'n_features': 16, 'epochs': 30, 'learning_rate': 0.05, 'lambda_reg': 0.01, 'n_hidden_units': 32, 'n_hidden_layers': 1}
Binary Neural Network - Empirical Risk Estimate: 0.43897262965985584, Accuracy: 0.7506, F1 score: 0.4066613515367429, Precision: 0.2672795367251504, Recall: 0.85771651065744, AUC: 0.8732469714613273
    Pred Neg  Pred Pos
Actual Neg      655      255
Actual Pos       13       77
...
...
Receiver Operating Characteristic - Binary Neural Network



The figure is a Receiver Operating Characteristic (ROC) plot titled "Receiver Operating Characteristic - Binary Neural Network". The x-axis is labeled "False Positive Rate" and ranges from 0.0 to 1.0. The y-axis is labeled "True Positive Rate" and ranges from 0.0 to 1.0. A solid orange line represents the "ROC curve (area = 0.87)". A dashed blue diagonal line represents a random classifier. The ROC curve starts at (0,0), rises steeply, and then levels off towards (1,1).


```

6. Results – Comparison

- Logistic Regression model slightly outperforms Neural Network model regarding AUC and F1
 - Logistic Regression model is better at separating classes across decision thresholds and at threshold of 0.5
- Neural Network model slightly outperforms Logistic Regression model in recall, but lacks behind in precision
 - It is more liberal: less false negatives, but more false positives

7. Conclusion

- The Logistic Regression model slightly outperforms the Neural Network
 - AUC metric is especially important in medical context
 - Minimizing false negatives may be more important than maximizing F1
- The competitive edge of the Logistic Regression model is not large
- Neural Network's ability to derive non-linear relations on its own does not grant performance edge, given the data set at hand

7. Conclusion

- Future work could focus on:
 - Leveraging more comprehensive data set and more complex Neural Network architecture: Might boost Neural Network performance
 - Creating higher-order polynomials and use dual decision functions: Might boost Logistic Regression performance and computational efficiency – but might as well lead to overfitting
 - Implementing adaptive learning in Neural Network: potentially allow convergence in better minima